

# Assessment of the Conditions of Wells in Tanzania

## Business Understanding

Tanzania, a developing country, has a problem of providing water to its fast growing population of 57 million. The country has already established wells that are expected to provide the much needed water, however, some of the wells require repairing for the goal of enough water supply to be met. Enthusiastic Environmentalists (EE), a renowned NGO is interested in locating the wells that require repairing and repair them, and enable the country curb the water problem. I am the data scientist tasked with developing a predictive model to know the wells that require repairing. This will enable them cut the cost of surveying, to pinpoint the exact wells that will require any repairing as using the model is easier to identify the wells without having to spend on nation-wide physical assessment.

## Data Understanding

```
In [1]: # importing libraries and functions
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn
from sklearn.preprocessing import MinMaxScaler, OneHotEncoder
from sklearn.metrics import accuracy_score,
classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from imblearn.over_sampling import SMOTE
```

```
In [2]: # Setting the display of columns to a maximum of 500
pd.set_option('display.max_columns', 500)
```

```
In [3]: # reading the two data files
X = pd.read_csv('Independent_Variables.csv')
y = pd.read_csv('Dependent_Variable.csv')
```

X dataframe is the dataset containing all the factors that relate to the target variable. The y dataframe contains the target variable which is the condition of the well.

```
In [4]: # viewing the first 5 rows of the independent variables to get a preview
X.head()
```

```
Out[4]:
```

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	none

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt_name
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Zahanati
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	Kwa Mahundi
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Zahanati Ya Nanyumbu
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Shuleni

In [5]:

```
# Viewing the first 5 rows of the dependent variable to get a preview
y.head()
```

Out[5]:

	id	status_group
0	69572	functional
1	8776	functional
2	34310	functional
3	67743	non functional
4	19728	functional

In [6]:

```
# Merging the two data sets for easier cleaning and manipulation
df = pd.concat([y['status_group'],X],axis = 1)
df.head()
```

Out[6]:

	status_group	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude
0	functional	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.85632
1	functional	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.14746
2	functional	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.82132
3	non functional	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.15529

	status_group	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude
4	functional	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.82535

In [7]:

```
# to get a general understanding of our data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   status_group                          59400 non-null  object
1   id                                    59400 non-null  int64
2   amount_tsh                           59400 non-null  float64
3   date_recorded                         59400 non-null  object
4   funder                                55765 non-null  object
5   gps_height                           59400 non-null  int64
6   installer                             55745 non-null  object
7   longitude                             59400 non-null  float64
8   latitude                             59400 non-null  float64
9   wpt_name                              59400 non-null  object
10  num_private                           59400 non-null  int64
11  basin                                 59400 non-null  object
12  subvillage                           59029 non-null  object
13  region                               59400 non-null  object
14  region_code                           59400 non-null  int64
15  district_code                         59400 non-null  int64
16  lga                                   59400 non-null  object
17  ward                                  59400 non-null  object
18  population                            59400 non-null  int64
19  public_meeting                       56066 non-null  object
20  recorded_by                           59400 non-null  object
21  scheme_management                     55523 non-null  object
22  scheme_name                           31234 non-null  object
23  permit                                56344 non-null  object
24  construction_year                     59400 non-null  int64
25  extraction_type                       59400 non-null  object
26  extraction_type_group                  59400 non-null  object
27  extraction_type_class                  59400 non-null  object
28  management                            59400 non-null  object
29  management_group                      59400 non-null  object
30  payment                               59400 non-null  object
31  payment_type                          59400 non-null  object
32  water_quality                         59400 non-null  object
33  quality_group                         59400 non-null  object
34  quantity                             59400 non-null  object
35  quantity_group                        59400 non-null  object
36  source                               59400 non-null  object
37  source_type                           59400 non-null  object
38  source_class                          59400 non-null  object
39  waterpoint_type                       59400 non-null  object
40  waterpoint_type_group                  59400 non-null  object
dtypes: float64(3), int64(7), object(31)
memory usage: 18.6+ MB
```

In [8]:

```
# to get descriptive statistics of our data
df.describe()
```

Out[8]:

	id	amount_tsh	gps_height	longitude	latitude	num_private	region_cc
<b>count</b>	59400.000000	59400.000000	59400.000000	59400.000000	5.940000e+04	59400.000000	59400.000000
<b>mean</b>	37115.131768	317.650385	668.297239	34.077427	-5.706033e+00	0.474141	15.297000
<b>std</b>	21453.128371	2997.574558	693.116350	6.567432	2.946019e+00	12.236230	17.587400
<b>min</b>	0.000000	0.000000	-90.000000	0.000000	-1.164944e+01	0.000000	1.000000
<b>25%</b>	18519.750000	0.000000	0.000000	33.090347	-8.540621e+00	0.000000	5.000000
<b>50%</b>	37061.500000	0.000000	369.000000	34.908743	-5.021597e+00	0.000000	12.000000
<b>75%</b>	55656.500000	20.000000	1319.250000	37.178387	-3.326156e+00	0.000000	17.000000
<b>max</b>	74247.000000	350000.000000	2770.000000	40.345193	-2.000000e-08	1776.000000	99.000000

In [9]:

```
# Trying to showcase the null values in the dataset
df.isna().sum()
```

```
Out[9]: status_group      0
        id                0
        amount_tsh        0
        date_recorded     0
        funder            3635
        gps_height         0
        installer         3655
        longitude         0
        latitude          0
        wpt_name          0
        num_private       0
        basin             0
        subvillage        371
        region            0
        region_code       0
        district_code     0
        lga               0
        ward              0
        population        0
        public_meeting    3334
        recorded_by       0
        scheme_management 3877
        scheme_name       28166
        permit            3056
        construction_year 0
        extraction_type    0
        extraction_type_group 0
        extraction_type_class 0
        management        0
        management_group  0
        payment            0
        payment_type       0
        water_quality      0
        quality_group      0
        quantity           0
        quantity_group     0
        source             0
        source_type        0
        source_class       0
        waterpoint_type    0
        waterpoint_type_group 0
        dtype: int64
```

```
In [10]: # Investigate the ternary target variable
         df['status_group'].value_counts()
```

```
Out[10]: functional      32259
         non functional   22824
         functional needs repair  4317
         Name: status_group, dtype: int64
```

```
In [11]: # Turn the target variable into binary
         df['status_group'] =
         df['status_group'].replace({"functional": "No_repair",
                                     "non functional": "Repair",
                                     "functional needs repair": "Repair"})
         df['status_group'].value_counts()
```

```
Out[11]: No_repair      32259
         Repair        27141
         Name: status_group, dtype: int64
```

From the data provided we see that 32,259 wells do not need any repair while 27,141 wells need to be repaired. We therefore use the various factors that predict whether a well would require repairing or not.

```
In [12]: # I chose to drop all the columns that had null values  
df = df.dropna(axis=1)  
df.isna().sum()
```

```
Out[12]: status_group      0  
id      0  
amount_tsh      0  
date_recorded      0  
gps_height      0  
longitude      0  
latitude      0  
wpt_name      0  
num_private      0  
basin      0  
region      0  
region_code      0  
district_code      0  
lga      0  
ward      0  
population      0  
recorded_by      0  
construction_year      0  
extraction_type      0  
extraction_type_group      0  
extraction_type_class      0  
management      0  
management_group      0  
payment      0  
payment_type      0  
water_quality      0  
quality_group      0  
quantity      0  
quantity_group      0  
source      0  
source_type      0  
source_class      0  
waterpoint_type      0  
waterpoint_type_group      0  
dtype: int64
```

```
In [13]: # Understanding the shape of the dataframe  
print(f''' The data has {df.shape[0]} rows and {df.shape[1]} columns''')
```

The data has 59400 rows and 34 columns

```
In [14]: # showing the columns in the dataframe  
df.columns
```

```
Out[14]: Index(['status_group', 'id', 'amount_tsh', 'date_recorded', 'gps_height',  
              'longitude', 'latitude', 'wpt_name', 'num_private', 'basin', 'region',  
              'region_code', 'district_code', 'lga', 'ward', 'population',  
              'recorded_by', 'construction_year', 'extraction_type',  
              'extraction_type_group', 'extraction_type_class', 'management',  
              'management_group', 'payment', 'payment_type', 'water_quality',  
              'quality_group', 'quantity', 'quantity_group', 'source', 'source_type',  
              'source_class', 'waterpoint_type', 'waterpoint_type_group'],  
             dtype='object')
```

```
In [15]: # Creating the dependent and independent variables  
X_ = df.drop(columns=['status_group','id'])  
y_ = df[['status_group']]
```

## Data Preparation

After reading the two datasets, getting to understand the content in both and cleaning the data to make it usable for modelling, we now head into preparing the data through transformations for the purposes of modelling.

```
In [16]: # splitting the dataset into training and testing datasets  
X_train,X_test,y_train,y_test = train_test_split(X_,y_,random_state=42)
```

```
In [17]: # One hot encoding the y variable  
ohe = OneHotEncoder(handle_unknown='ignore')  
ohe.fit(y_train)  
y_train_ = ohe.transform(y_train).toarray()  
y_train_enc =  
pd.DataFrame(y_train_,columns=ohe.get_feature_names(y_train.columns))  
y_train_enc = y_train_enc[['status_group_Repair']]  
y_train_enc
```

Out[17]:

	status_group_Repair
0	1.0
1	0.0
2	0.0
3	0.0
4	1.0
...	...
44545	0.0
44546	0.0
44547	1.0
44548	0.0
44549	1.0

44550 rows × 1 columns

In [18]:

```
#Onehotencoding the categorical variables in the independent variables dataset
X_train_categorical = X_train.drop(columns=
['date_recorded', 'wpt_name']).select_dtypes(include=['object'])
ohe.fit(X_train_categorical)
X_train_ = ohe.transform(X_train_categorical).toarray()
X_train_enc =
pd.DataFrame(X_train_, columns=ohe.get_feature_names(X_train_categorical.columns))
X_train_enc
```

Out[18]:

	basin_Internal	basin_Lake Nyasa	basin_Lake Rukwa	basin_Lake Tanganyika	basin_Lake Victoria	basin_Pangani	basin_Rufiji	basin_ /
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	0.0	1.0	0.0	0.0	
2	0.0	0.0	0.0	0.0	1.0	0.0	0.0	
3	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
4	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
...	...	...	...	...	...	...	...	
44545	0.0	0.0	0.0	0.0	0.0	0.0	1.0	
44546	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
44547	0.0	0.0	0.0	0.0	0.0	1.0	0.0	
44548	0.0	0.0	0.0	0.0	1.0	0.0	0.0	
44549	0.0	0.0	0.0	0.0	0.0	1.0	0.0	



44550 rows × 2353 columns

```
In [19]: # Creating a variable with numeric variables
X_train_num = X_train.select_dtypes(exclude=['object'])
X_train_num = X_train_num.reset_index()
```

```
In [20]: # Scaling the numeric variables
scaler = MinMaxScaler()
scaler.fit(X_train_num.drop(columns='index'))
X_train_scaled = scaler.transform(X_train_num.drop(columns='index'))
X_train_scaled_df =
pd.DataFrame(X_train_scaled, columns=X_train_num.drop(columns='index').columns)
X_train_scaled_df
```

```
Out[20]:
```

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population
0	0.000057	0.138722	0.944941	0.477474	0.0	0.051020	0.0125	0.002623
1	0.000000	0.022238	0.000000	1.000000	0.0	0.163265	0.0125	0.000000
2	0.000000	0.022238	0.825683	0.758435	0.0	0.183673	0.0500	0.000000
3	0.000000	0.566537	0.862136	0.584350	0.0	0.122449	0.0500	0.000754
4	0.000000	0.206848	0.859110	0.080872	0.0	0.091837	0.0375	0.000033
...	...	...	...	...	...	...	...	...
44545	0.002857	0.137663	0.901399	0.246765	0.0	0.040816	0.0500	0.008361
44546	0.002857	0.637487	0.855902	0.161367	0.0	0.102041	0.0625	0.001148
44547	0.000000	0.017649	0.966024	0.534671	0.0	0.030612	0.0625	0.032781
44548	0.000000	0.022238	0.850574	0.733278	0.0	0.163265	0.0750	0.000000
44549	0.000000	0.477586	0.932612	0.724325	0.0	0.020408	0.0125	0.000033

44550 rows × 9 columns

```
In [21]: # Creating one dataset with all independent variables ready for
modelling
X_values_train = pd.concat([X_train_scaled_df, X_train_enc], axis=1)
X_values_train
```

```
Out[21]:
```

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population
0	0.000057	0.138722	0.944941	0.477474	0.0	0.051020	0.0125	0.002623
1	0.000000	0.022238	0.000000	1.000000	0.0	0.163265	0.0125	0.000000
2	0.000000	0.022238	0.825683	0.758435	0.0	0.183673	0.0500	0.000000
3	0.000000	0.566537	0.862136	0.584350	0.0	0.122449	0.0500	0.000754
4	0.000000	0.206848	0.859110	0.080872	0.0	0.091837	0.0375	0.000033

	amount_tsh	gps_height	longitude	latitude	num_private	region_code	district_code	population
...	...	...	...	...	...	...	...	...
44545	0.002857	0.137663	0.901399	0.246765	0.0	0.040816	0.0500	0.00836
44546	0.002857	0.637487	0.855902	0.161367	0.0	0.102041	0.0625	0.001148
44547	0.000000	0.017649	0.966024	0.534671	0.0	0.030612	0.0625	0.03278
44548	0.000000	0.022238	0.850574	0.733278	0.0	0.163265	0.0750	0.000000
44549	0.000000	0.477586	0.932612	0.724325	0.0	0.020408	0.0125	0.00003

44550 rows × 2362 columns

## Modeling and Evaluation

After processing the data and getting it ready for modelling, we head into creating predictive models from which the model with the best performance will be selected and used by the NGO in making predictions.

```
In [22]: # Creating the baseline model which is a Logistic regression
baseline_model = LogisticRegression(max_iter=1000)
baseline_model.fit(X=X_values_train,y=np.ravel(y_train_enc))
```

```
Out[22]: LogisticRegression(max_iter=1000)
```

```
In [23]: # Predicting the y values with train x values
y_pred = baseline_model.predict(X_values_train)
```

```
In [24]: # Seeing the performance of the model with the train dataset
train_report = classification_report(y_true=y_train_enc,y_pred=y_pred)
print(train_report)
print(f'''{accuracy_score(y_true=y_train_enc,y_pred=y_pred)}''')
```

	precision	recall	f1-score	support
0.0	0.79	0.87	0.83	24161
1.0	0.82	0.73	0.77	20389
accuracy			0.80	44550
macro avg	0.81	0.80	0.80	44550
weighted avg	0.81	0.80	0.80	44550

0.804287317620651

```
In [25]: # Preprocessing the test dataset
ohe.fit(y_train)
y_test_ = ohe.transform(y_test).toarray()
y_test_enc =
pd.DataFrame(y_test_,columns=ohe.get_feature_names(y_test.columns))
```

```

X_test_categorical = X_test.drop(columns=
['date_recorded', 'wpt_name']).select_dtypes(include=['object'])
ohe.fit(X_train_categorical)
X_test_ = ohe.transform(X_test_categorical).toarray()
X_test_enc =
pd.DataFrame(X_test_, columns=ohe.get_feature_names(X_test_categorical.colu
X_test_num = X_test.select_dtypes(exclude=['object'])
X_test_num = X_test_num.reset_index()
X_test_scaled = scaler.transform(X_test_num.drop(columns='index'))
X_test_scaled_df =
pd.DataFrame(X_test_scaled, columns=X_test_num.drop(columns='index').column
X_test_scaled_df
X_values_test = pd.concat([X_test_scaled_df, X_test_enc], axis=1)

```

```

In [26]: # predicting the y values using the test x values
y_pred2 = baseline_model.predict(X_values_test)

```

```

In [27]: # creating the y variable
y_test_enc = y_test_enc[['status_group_Repair']]

```

```

In [28]: # Seeing the performance of the model using the test dataset
test_report = classification_report(y_true=y_test_enc, y_pred=y_pred2)
print(test_report)


```

	precision	recall	f1-score	support
0.0	0.78	0.85	0.81	8098
1.0	0.79	0.71	0.75	6752
accuracy			0.78	14850
macro avg	0.79	0.78	0.78	14850
weighted avg	0.79	0.78	0.78	14850

```

In [29]: # Visualizing the confusion matrix
sns.heatmap(confusion_matrix(y_true = y_test_enc, y_pred=y_pred2), annot =
True);

```



```

In [30]: # Using the AUC score to test the performance of the model

```

```
sklearn.metrics.roc_auc_score(y_true=y_test_enc,y_score=y_pred2)
```

Out[30]: 0.7784189919048529

```
In [31]: # Creating a decision tree model
improved_model = sklearn.tree.DecisionTreeClassifier()
improved_model.fit(X=X_values_train,y=y_train_enc)
```

Out[31]: DecisionTreeClassifier()

```
In [32]: # predicting y values using the decision tree model
y_pred3 = improved_model.predict(X_values_test)
```

```
In [33]: # testing the performance of the model using the AUC score
sklearn.metrics.roc_auc_score(y_true=y_test_enc,y_score=y_pred3)
```

Out[33]: 0.7842471453076589

```
In [34]: # classification report
test_report_improved =
classification_report(y_true=y_test_enc,y_pred=y_pred3)
print(test_report_improved)
```

	precision	recall	f1-score	support
0.0	0.80	0.81	0.81	8098
1.0	0.77	0.76	0.76	6752
accuracy			0.79	14850
macro avg	0.79	0.78	0.78	14850
weighted avg	0.79	0.79	0.79	14850

## Conclusion

For this model the selected metric used to determine the most appropriate model was AUC score since it is best for binary target variables. The decision tree classification model (improved\_model) had an AUC score of 78%, indicating that it is the best model to use. The NGO should then use the improved model to identify wells that need to be repaired.